The dominant sequence transduction models are based on complex recurrent neural networks that inc
Attention mechanisms have become an integral part of compelling sequence modeling and transduc-\n
oder inserts a third sub-layer, which performs multi-head\nattention over the output of the encoder sta
In this work we employ h= 8 parallel attention layers, or heads. For each of these we use h=dv=dmode
In this section we compare various aspects of self-attention layers to the recurrent and convolu-\ntiona
MT 2014 English-German dataset consisting of about 4.5 million sentences. Sentences were encoded
The Transformer (big) model trained for English-to-French used a lower drop rate than the previous st
We trained a 4-layer transformer with dmodel = 1024 on the Wall Street Journal (WSJ) portion of the P
Arxiv preprint arXiv:1601.06733 , 2016.\n[1] Kyunghyun Cho, Bart van Merrienboer, Caglar Gulcehre,
A structured self-attentive sentence embedding. Factorization tricks for LSTM networks. Multi-task seq
Rethinking the inception architecture for computer vision. Grammar as a foreign language. Google’s n
The attention heads exhibit behaviour that seems related to the structure of the\nsentence. We give tw